

clean__eda

January 26, 2026

1 Exploratory Data Analysis: Delhi NCR Air Quality Dataset

MSE 546 Project - AQI Forecasting (6-hour ahead prediction)

Date: January 26, 2026

Objective: Understand the data characteristics and patterns to inform baseline model development and future feature engineering.

1.1 1. Setup and Data Loading

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime

# Set style for better visualizations
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")
%matplotlib inline

# Display settings
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 100)
pd.set_option('display.precision', 2)

[2]: # Load data
PATH = "delhi_ncr_aqi_dataset.csv"
df = pd.read_csv(PATH)

print(f"Dataset Shape: {df.shape[0]:,} rows × {df.shape[1]} columns")
print(f"\nColumns: {df.columns.tolist()}")
```

Dataset Shape: 201,664 rows × 25 columns

Columns: ['datetime', 'date', 'year', 'month', 'day', 'hour', 'day_of_week', 'is_weekend', 'season', 'city', 'station', 'latitude', 'longitude', 'pm25',

```
'pm10', 'no2', 'so2', 'co', 'o3', 'temperature', 'humidity', 'wind_speed',  
'visibility', 'aqi', 'aqi_category']
```

1.2 2. Data Preprocessing

```
[3]: # Parse datetime  
if "datetime" in df.columns:  
    df["datetime"] = pd.to_datetime(df["datetime"], errors="coerce")  
elif "date" in df.columns and "hour" in df.columns:  
    df["datetime"] = pd.to_datetime(df["date"], errors="coerce") + pd.  
    ↪to_timedelta(df["hour"], unit="h")  
  
# Add day of week if not present  
if 'day_of_week' not in df.columns:  
    df['day_of_week'] = df['datetime'].dt.dayofweek  
  
# Ensure numeric columns are numeric  
num_cols = ["pm25", "pm10", "no2", "so2", "co", "o3",  
            "temperature", "humidity", "wind_speed", "visibility", "aqi",  
            "latitude", "longitude", "year", "month", "day", "hour",  
            ↪"is_weekend"]  
  
for c in num_cols:  
    if c in df.columns:  
        df[c] = pd.to_numeric(df[c], errors="coerce")  
  
# Sort by time  
df = df.sort_values(["station", "datetime"]).reset_index(drop=True)  
  
print(" Data preprocessing complete")
```

Data preprocessing complete

1.3 3. Data Quality Overview

```
[4]: # Check for duplicates  
duplicates = df.duplicated().sum()  
print(f"Duplicate rows: {duplicates}")  
  
# Missing values  
missing = df.isna().sum()  
missing_pct = (missing / len(df) * 100).round(2)  
missing_df = pd.DataFrame({  
    'Missing Count': missing,  
    'Missing %': missing_pct  
}).sort_values('Missing Count', ascending=False)  
  
print("\nMissing Values:")
```

```
print(missing_df[missing_df['Missing Count'] > 0])

if missing_df['Missing Count'].sum() == 0:
    print(" No missing values in the dataset!")
```

Duplicate rows: 0

Missing Values:

Empty DataFrame

Columns: [Missing Count, Missing %]

Index: []

No missing values in the dataset!

```
[5]: # Data types and basic info
print("\nData Types:")
print(df.dtypes)

print("\n" + "="*50)
print("First few rows:")
df.head()
```

Data Types:

datetime	datetime64[ns]
date	object
year	int64
month	int64
day	int64
hour	int64
day_of_week	object
is_weekend	int64
season	object
city	object
station	object
latitude	float64
longitude	float64
pm25	float64
pm10	float64
no2	float64
so2	float64
co	float64
o3	float64
temperature	float64
humidity	int64
wind_speed	float64
visibility	float64
aqi	int64
aqi_category	object

dtype: object

=====

First few rows:

```
[5]:
```

	datetime	date	year	month	day	hour	day_of_week	\
0	2020-01-01 06:00:00	2020-01-01	2020	1	1	6	Wednesday	
1	2020-01-01 12:00:00	2020-01-01	2020	1	1	12	Wednesday	
2	2020-01-01 18:00:00	2020-01-01	2020	1	1	18	Wednesday	
3	2020-01-01 23:00:00	2020-01-01	2020	1	1	23	Wednesday	
4	2020-01-02 06:00:00	2020-01-02	2020	1	2	6	Thursday	

	is_weekend	season	city	station	latitude	longitude	pm25	\
0	0	winter	Delhi	Anand Vihar, Delhi	28.65	77.32	371.8	
1	0	winter	Delhi	Anand Vihar, Delhi	28.65	77.32	301.1	
2	0	winter	Delhi	Anand Vihar, Delhi	28.65	77.32	334.0	
3	0	winter	Delhi	Anand Vihar, Delhi	28.65	77.32	403.8	
4	0	winter	Delhi	Anand Vihar, Delhi	28.65	77.32	448.4	

	pm10	no2	so2	co	o3	temperature	humidity	wind_speed	\
0	739.4	119.6	47.7	5.19	12.3	9.4	100	3.6	
1	588.8	117.9	39.3	4.32	15.8	20.6	50	5.9	
2	602.6	150.1	36.3	7.13	14.3	12.4	56	4.5	
3	841.8	142.0	30.3	4.90	13.2	14.4	48	5.8	
4	749.2	134.8	39.9	7.62	15.1	10.0	98	3.3	

	visibility	aqi	aqi_category
0	1.2	500	Severe
1	1.4	500	Severe
2	1.1	500	Severe
3	1.4	500	Severe
4	0.5	500	Severe

1.4 4. Temporal Coverage

Key Finding: Data has only 4 measurements per day (hours: 6, 12, 18, 23)

```
[7]: # Time range
print(f"Time Range: {df['datetime'].min()} to {df['datetime'].max()}")
print(f"Duration: {(df['datetime'].max() - df['datetime'].min()).days} days\n")

# Hours available per day
hour_counts = df['hour'].value_counts().sort_index()
print("Observations per hour of day:")
print(hour_counts)
```

Time Range: 2020-01-01 06:00:00 to 2025-12-31 23:00:00

Duration: 2191 days

Observations per hour of day:

hour

6 50416

12 50416

18 50416

23 50416

Name: count, dtype: int64

1.5 5. Spatial Coverage

```
[8]: # Cities and stations
print("Cities in dataset:")
city_counts = df['city'].value_counts()
print(city_counts)

print(f"\nTotal unique stations: {df['station'].nunique()}")

# Stations per city
stations_per_city = df.groupby('city')['station'].nunique().
    ↪sort_values(ascending=False)
print("\nStations per city:")
print(stations_per_city)
```

Cities in dataset:

city

Delhi 122752

Noida 26304

Faridabad 17536

Ghaziabad 17536

Gurugram 17536

Name: count, dtype: int64

Total unique stations: 23

Stations per city:

city

Delhi 14

Noida 3

Faridabad 2

Ghaziabad 2

Gurugram 2

Name: station, dtype: int64

```
[9]: # List all stations by city
print("\nAll stations by city:")
print("="*60)
for city in df['city'].unique():
    stations = df[df['city'] == city]['station'].unique()
```

```

print(f"\n{city} ({len(stations)} stations):")
for i, station in enumerate(sorted(stations), 1):
    print(f"  {i}. {station}")

```

All stations by city:

=====

Delhi (14 stations):

1. Anand Vihar, Delhi
2. Bawana, Delhi
3. Dwarka Sec 8, Delhi
4. ITO, Delhi
5. Jahangirpuri, Delhi
6. Mandir Marg, Delhi
7. NSIT Dwarka, Delhi
8. Okhla Phase 2, Delhi
9. Punjabi Bagh, Delhi
10. RK Puram, Delhi
11. Rohini, Delhi
12. Shadipur, Delhi
13. Siri Fort, Delhi
14. Wazirpur, Delhi

Faridabad (2 stations):

1. Faridabad New Town
2. Faridabad Sec 16A

Ghaziabad (2 stations):

1. Ghaziabad Loni
2. Ghaziabad Vasundhara

Noida (3 stations):

1. Greater Noida
2. Noida Sec 125
3. Noida Sec 62

Gurugram (2 stations):

1. Gurugram Sec 51
2. Gurugram Vikas Sadan

```

[11]: # Visualize station locations
station_info = df.groupby('station').agg({
    'latitude': 'first',
    'longitude': 'first',
    'city': 'first',
    'aqi': 'mean'

```

```

}).reset_index()

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

# Plot 1: Station locations
for city in station_info['city'].unique():
    city_data = station_info[station_info['city'] == city]
    ax1.scatter(city_data['longitude'], city_data['latitude'],
                label=city, s=100, alpha=0.6)
ax1.set_xlabel('Longitude')
ax1.set_ylabel('Latitude')
ax1.set_title('Station Locations by City')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Station locations sized by mean AQI
scatter = ax2.scatter(station_info['longitude'], station_info['latitude'],
                      c=station_info['aqi'], s=station_info['aqi'],
                      alpha=0.6, cmap='YlOrRd')
ax2.set_xlabel('Longitude')
ax2.set_ylabel('Latitude')
ax2.set_title('Station Locations (size & color = mean AQI)')
plt.colorbar(scatter, ax=ax2, label='Mean AQI')
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



1.6 6. Target Variable: AQI Distribution

```
[14]: # Summary statistics for AQI
print("AQI Summary Statistics:")
print(df['aqi'].describe())
```

```
AQI Summary Statistics:
count      201664.00
mean        265.83
std         171.51
min          25.00
25%         103.00
50%         232.00
75%         464.00
max          500.00
Name: aqi, dtype: float64
```

```
[13]: # AQI categories
print("\nAQI Category Distribution:")
category_counts = df['aqi_category'].value_counts()
print(category_counts)
print(f"\nCategory percentages:")
print((category_counts / len(df) * 100).round(2))
```

```
AQI Category Distribution:
aqi_category
Severe      60102
Moderate    45397
Satisfactory 32860
Very Poor   29674
Poor        18843
Good        14788
Name: count, dtype: int64
```

```
Category percentages:
aqi_category
Severe      29.80
Moderate    22.51
Satisfactory 16.29
Very Poor   14.71
Poor         9.34
Good         7.33
Name: count, dtype: float64
```

```
[15]: # Visualize AQI distribution
fig, axes = plt.subplots(2, 2, figsize=(15, 10))
```



```

# Histogram
axes[0, 0].hist(df['aqi'], bins=50, edgecolor='black', alpha=0.7)
axes[0, 0].set_xlabel('AQI')
axes[0, 0].set_ylabel('Frequency')
axes[0, 0].set_title('AQI Distribution')
axes[0, 0].axvline(df['aqi'].mean(), color='red', linestyle='--', label=f'Mean: {df["aqi"].mean():.1f}')
axes[0, 0].axvline(df['aqi'].median(), color='green', linestyle='--', label=f'Median: {df["aqi"].median():.1f}')
axes[0, 0].legend()


# Box plot
axes[0, 1].boxplot(df['aqi'].dropna(), vert=True)
axes[0, 1].set_ylabel('AQI')
axes[0, 1].set_title('AQI Box Plot')
axes[0, 1].grid(True, alpha=0.3)

# Category bar plot
category_order = ['Good', 'Satisfactory', 'Moderate', 'Poor', 'Very Poor', 'Severe']
cat_counts = df['aqi_category'].value_counts().reindex(category_order, fill_value=0)
colors = ['green', 'lightgreen', 'yellow', 'orange', 'red', 'darkred']
axes[1, 0].bar(range(len(cat_counts)), cat_counts.values, color=colors, alpha=0.7)
axes[1, 0].set_xticks(range(len(cat_counts)))
axes[1, 0].set_xticklabels(cat_counts.index, rotation=45, ha='right')
axes[1, 0].set_ylabel('Count')
axes[1, 0].set_title('AQI Category Distribution')
axes[1, 0].grid(True, alpha=0.3, axis='y')

# Log-scale histogram for better view of tails
axes[1, 1].hist(df['aqi'], bins=50, edgecolor='black', alpha=0.7)
axes[1, 1].set_xlabel('AQI')
axes[1, 1].set_ylabel('Frequency (log scale)')
axes[1, 1].set_title('AQI Distribution (Log Scale)')
axes[1, 1].set_yscale('log')

plt.tight_layout()
plt.show()

```



1.7 7. Pollutant Distributions

```
[16]: # Pollutant summary statistics
pollutants = ['pm25', 'pm10', 'no2', 'so2', 'co', 'o3']
pollutant_stats = df[pollutants].describe().T
print("Pollutant Summary Statistics:")
print(pollutant_stats)
```

Pollutant Summary Statistics:

	count	mean	std	min	25%	50%	75%	max
pm25	201664.0	183.42	193.14	15.0	55.30	99.50	254.70	900.00
pm10	201664.0	348.57	370.02	24.0	104.10	189.80	481.20	1979.70
no2	201664.0	69.76	75.97	8.0	19.90	38.30	94.00	593.50
so2	201664.0	16.03	17.20	4.0	4.50	8.70	20.90	121.60
co	201664.0	3.03	3.28	0.3	0.87	1.69	4.12	22.67
o3	201664.0	27.19	13.57	12.0	18.00	23.30	31.80	84.00

```
[17]: # Visualize pollutant distributions
fig, axes = plt.subplots(2, 3, figsize=(18, 10))
axes = axes.flatten()


for i, pollutant in enumerate(pollutants):
    axes[i].hist(df[pollutant].dropna(), bins=50, edgecolor='black', alpha=0.7,
                color='steelblue')
```

```

axes[i].set_xlabel(pollutant.upper())
axes[i].set_ylabel('Frequency')
axes[i].set_title(f'{pollutant.upper()} Distribution')
axes[i].axvline(df[pollutant].mean(), color='red', linestyle='--',
                label=f'Mean: {df[pollutant].mean():.2f}')
axes[i].legend()
axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



1.8 8. Weather Variables

```

[18]: # Weather summary statistics
weather_vars = ['temperature', 'humidity', 'wind_speed', 'visibility']
weather_stats = df[weather_vars].describe().T
print("Weather Variable Summary Statistics:")
print(weather_stats)

```


Weather Variable Summary Statistics:

	count	mean	std	min	25%	50%	75%	max
temperature	201664.0	26.90	8.98	6.0	19.9	27.6	33.5	44.8
humidity	201664.0	61.04	24.74	20.0	39.0	58.0	83.0	100.0
wind_speed	201664.0	10.53	6.05	2.1	5.9	9.2	14.2	31.5
visibility	201664.0	5.01	3.69	0.2	1.9	4.4	7.1	15.0

```
[19]: # Visualize weather distributions
fig, axes = plt.subplots(2, 2, figsize=(15, 10))
axes = axes.flatten()

for i, var in enumerate(weather_vars):
    axes[i].hist(df[var].dropna(), bins=50, edgecolor='black', alpha=0.7,
    color='coral')
    axes[i].set_xlabel(var.replace('_', ' ').title())
    axes[i].set_ylabel('Frequency')
    axes[i].set_title(f'{var.replace("_", " ").title()} Distribution')
    axes[i].axvline(df[var].mean(), color='blue', linestyle='--',
    label=f'Mean: {df[var].mean():.2f}')
    axes[i].legend()
    axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



1.9 9. Correlation Analysis


```
[20]: # Correlation with AQI
feature_cols = pollutants + weather_vars
correlations = df[feature_cols + ['aqi']].corr()['aqi'].drop('aqi').
    ↪sort_values(ascending=False)

print("Correlation with AQI:")
print(correlations)

# Visualize
fig, ax = plt.subplots(figsize=(10, 6))
colors = ['green' if x > 0 else 'red' for x in correlations.values]
ax.barh(range(len(correlations)), correlations.values, color=colors, alpha=0.7)
ax.set_yticks(range(len(correlations)))
ax.set_yticklabels(correlations.index)
ax.set_xlabel('Correlation with AQI')
ax.set_title('Feature Correlations with AQI')
ax.axvline(0, color='black', linewidth=0.5)
ax.grid(True, alpha=0.3, axis='x')
plt.tight_layout()
plt.show()

print("\n Key Insights:")
print(f"    • Strongest positive: {correlations.idxmax()} ({correlations.max():.
    ↪3f})")
print(f"    • Strongest negative: {correlations.idxmin()} ({correlations.min():.
    ↪3f})")
```

```
Correlation with AQI:
pm25      0.80
pm10      0.80
co        0.78
no2       0.77
so2       0.76
humidity  -0.22
o3        -0.28
wind_speed -0.54
temperature -0.73
visibility -0.86
Name: aqi, dtype: float64
```



Key Insights:


- Strongest positive: pm25 (0.801)
- Strongest negative: visibility (-0.857)

```
[21]: # Full correlation matrix
corr_matrix = df[feature_cols + ['aqi']].corr()

fig, ax = plt.subplots(figsize=(12, 10))
im = ax.imshow(corr_matrix, cmap='coolwarm', vmin=-1, vmax=1, aspect='auto')
ax.set_xticks(range(len(corr_matrix.columns)))
ax.set_yticks(range(len(corr_matrix.columns)))
ax.set_xticklabels(corr_matrix.columns, rotation=45, ha='right')
ax.set_yticklabels(corr_matrix.columns)
ax.set_title('Correlation Matrix: Pollutants, Weather & AQI', fontsize=14,
             pad=20)

# Add correlation values
for i in range(len(corr_matrix)):
    for j in range(len(corr_matrix)):
        text = ax.text(j, i, f'{corr_matrix.iloc[i, j]:.2f}',
                       ha="center", va="center", color="black", fontsize=8)


plt.colorbar(im, ax=ax, label='Correlation')
plt.tight_layout()
plt.show()
```



1.10 10. Temporal Patterns

```
[22]: # AQI over time (daily mean)
daily_aqi = df.set_index('datetime')['aqi'].resample('D').mean()

fig, ax = plt.subplots(figsize=(15, 6))
ax.plot(daily_aqi.index, daily_aqi.values, alpha=0.7, linewidth=0.8)
ax.set_xlabel('Date')
ax.set_ylabel('Daily Mean AQI')
ax.set_title('Daily Mean AQI Over Time (All Stations)')
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




```
[23]: # AQI by hour of day
hourly_stats = df.groupby('hour')['aqi'].agg(['mean', 'std', 'median'])
print("AQI Statistics by Hour of Day:")
print(hourly_stats)

fig, ax = plt.subplots(figsize=(10, 6))
hours = hourly_stats.index
ax.plot(hours, hourly_stats['mean'], marker='o', label='Mean', linewidth=2)
ax.plot(hours, hourly_stats['median'], marker='s', label='Median', linewidth=2)
ax.fill_between(hours,
                hourly_stats['mean'] - hourly_stats['std'],
                hourly_stats['mean'] + hourly_stats['std'],
                alpha=0.3, label='±1 Std Dev')
ax.set_xlabel('Hour of Day')
ax.set_ylabel('AQI')
ax.set_title('AQI by Hour of Day')
ax.legend()
ax.grid(True, alpha=0.3)
ax.set_xticks(hours)
plt.tight_layout()
plt.show()

print("\n Insight: Hour 18 (6 PM) has highest AQI, Hour 12 (noon) has lowest")
```

AQI Statistics by Hour of Day:


	mean	std	median
hour			
6	271.31	178.68	229.0
12	243.36	166.85	168.0
18	281.62	163.12	292.0
23	267.04	174.65	230.0



Insight: Hour 18 (6 PM) has highest AQI, Hour 12 (noon) has lowest

```
[24]: # AQI by day of week
dow_stats = df.groupby('day_of_week')['aqi'].agg(['mean', 'median'])
dow_names = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']


fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(range(7), dow_stats['mean'].values, alpha=0.7, color='steelblue',
      edgecolor='black')
ax.set_xticks(range(7))
ax.set_xticklabels(dow_names)
ax.set_xlabel('Day of Week')
ax.set_ylabel('Mean AQI')
ax.set_title('Mean AQI by Day of Week')
ax.grid(True, alpha=0.3, axis='y')
plt.tight_layout()
plt.show()
```



```
[25]: # AQI by month
monthly_stats = df.groupby('month')['aqi'].agg(['mean', 'median'])
month_names = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
               'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']

fig, ax = plt.subplots(figsize=(12, 6))
ax.plot(range(1, 13), monthly_stats['mean'].values, marker='o', linewidth=2,
        markersize=8)
ax.set_xticks(range(1, 13))
ax.set_xticklabels(month_names)
ax.set_xlabel('Month')
ax.set_ylabel('Mean AQI')
ax.set_title('Mean AQI by Month (Strong Seasonal Pattern)')
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print("\n Insight: Winter months (Nov-Jan) have much higher AQI")
```



Insight: Winter months (Nov-Jan) have much higher AQI

```
[46]: # AQI by season
season_stats = df.groupby('season')['aqi'].agg(['mean', 'median', 'std', 'count'])
season_order = ['winter', 'summer', 'monsoon', 'post_monsoon']
season_stats = season_stats.reindex(season_order)

print("AQI Statistics by Season:")
print(season_stats)


fig, ax = plt.subplots(figsize=(10, 6))
colors = ['blue', 'yellow', 'green', 'orange']
ax.bar(range(len(season_stats)), season_stats['mean'].values,
       alpha=0.7, color=colors, edgecolor='black')
ax.set_xticks(range(len(season_stats)))
ax.set_xticklabels(season_stats.index, rotation=0)
ax.set_xlabel('Season')
ax.set_ylabel('Mean AQI')
ax.set_title('Mean AQI by Season')
ax.grid(True, alpha=0.3, axis='y')

# Add value labels
for i, v in enumerate(season_stats['mean'].values):
    ax.text(i, v + 5, f'{v:.1f}', ha='center', fontweight='bold')

plt.tight_layout()
plt.show()
```

AQI Statistics by Season:

	mean	median	std	count
season				
winter	457.79	500.0	59.26	66424
summer	204.79	198.0	64.94	50784
monsoon	82.89	82.0	34.76	67344
post_monsoon	421.81	426.0	75.01	17112



1.11 11. Spatial Patterns


```
[47]: # AQI by city
city_stats = df.groupby('city')['aqi'].agg(['mean', 'median', 'std']).
    ↪sort_values('mean', ascending=False)
print("AQI Statistics by City:")
print(city_stats)

fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(range(len(city_stats)), city_stats['mean'].values, alpha=0.7,
    ↪edgecolor='black')
ax.set_xticks(range(len(city_stats)))
ax.set_xticklabels(city_stats.index, rotation=45, ha='right')
ax.set_xlabel('City')
ax.set_ylabel('Mean AQI')
ax.set_title('Mean AQI by City')
ax.grid(True, alpha=0.3, axis='y')
```

```
plt.tight_layout()
plt.show()
```

AQI Statistics by City:

	mean	median	std
city			
Ghaziabad	275.08	250.0	171.84
Delhi	268.10	238.0	171.52
Faridabad	263.27	225.5	171.45
Gurugram	258.22	214.0	171.23
Noida	255.87	209.0	170.83



```
[48]: # Top 20 stations by mean AQI
station_stats = df.groupby('station')['aqi'].agg(['mean', 'median', 'std',
↪ 'count']).sort_values('mean', ascending=False)
top_20_stations = station_stats.head(20)

print("Top 20 Stations by Mean AQI:")
print(top_20_stations)

fig, ax = plt.subplots(figsize=(12, 8))
ax.barh(range(len(top_20_stations)), top_20_stations['mean'].values, alpha=0.7,
↪ edgecolor='black')
ax.set_yticks(range(len(top_20_stations)))
ax.set_yticklabels(top_20_stations.index)
```


```

ax.set_xlabel('Mean AQI')
ax.set_title('Top 20 Stations by Mean AQI')
ax.grid(True, alpha=0.3, axis='x')
plt.tight_layout()
plt.show()

```

Top 20 Stations by Mean AQI:

	mean	median	std	count
station				
Anand Vihar, Delhi	295.17	298.0	170.34	8768
Jahangirpuri, Delhi	289.72	287.0	170.84	8768
Wazirpur, Delhi	286.59	277.0	171.10	8768
Bawana, Delhi	283.69	270.0	171.39	8768
Ghaziabad Loni	279.57	262.0	171.43	8768
ITO, Delhi	276.97	252.0	171.45	8768
Okhla Phase 2, Delhi	274.13	250.0	172.02	8768
Ghaziabad Vasundhara	270.59	240.0	172.14	8768
Punjabi Bagh, Delhi	270.28	239.0	171.19	8768
Shadipur, Delhi	266.13	230.0	171.63	8768
Faridabad Sec 16A	266.11	232.0	171.41	8768
Noida Sec 62	263.22	225.0	171.28	8768
Rohini, Delhi	263.04	225.0	171.58	8768
Gurugram Vikas Sadan	260.43	217.0	171.54	8768
Faridabad New Town	260.43	218.0	171.44	8768
Gurugram Sec 51	256.01	211.0	170.90	8768
Noida Sec 125	255.90	210.0	170.87	8768
RK Puram, Delhi	255.61	210.0	170.92	8768
Siri Fort, Delhi	251.16	204.0	170.13	8768
Greater Noida	248.49	196.0	170.04	8768



1.12 12. AQI Persistence Check (Important for Baseline)

Since we're predicting 6 hours ahead, let's check how much AQI changes over that time period.

```
[49]: # Create target variable (next measurement AQI)
df_sorted = df.sort_values(['station', 'datetime']).copy()
df_sorted['aqi_next'] = df_sorted.groupby('station')['aqi'].shift(-1)

# Remove rows where we can't compute next AQI
df_persist = df_sorted.dropna(subset=['aqi_next']).copy()

print(f"Rows with both current and next AQI: {len(df_persist):,}")
print(f"Rows dropped (last measurement per station): {len(df_sorted) - \
    ↪ len(df_persist):,}")
```

Rows with both current and next AQI: 201,641

Rows dropped (last measurement per station): 23

```
[50]: # Calculate AQI change
df_persist['aqi_change'] = df_persist['aqi_next'] - df_persist['aqi']
df_persist['aqi_pct_change'] = (df_persist['aqi_change'] / df_persist['aqi'] * \
    ↪ 100)

print("AQI Change Statistics (t to t+1):")
```

```
print(df_persist['aqi_change'].describe())
print("\nAQI % Change Statistics:")
print(df_persist['aqi_pct_change'].describe())
```

AQI Change Statistics (t to t+1):

```
count    2.02e+05
mean     -5.80e-04
std       5.58e+01
min      -2.97e+02
25%      -2.90e+01
50%       0.00e+00
75%       2.50e+01
max       3.43e+02
Name: aqi_change, dtype: float64
```

AQI % Change Statistics:

```
count    201641.00
mean         6.65
std         41.83
min        -81.73
25%        -15.88
50%         0.00
75%         15.58
max        794.74
Name: aqi_pct_change, dtype: float64
```

```
[51]: # Scatter plot: current AQI vs next AQI
sample = df_persist.sample(min(5000, len(df_persist)), random_state=42)

fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Scatter plot
axes[0].scatter(sample['aqi'], sample['aqi_next'], alpha=0.3, s=10)
axes[0].plot([0, 500], [0, 500], 'r--', label='Perfect persistence',
             linewidth=2)
axes[0].set_xlabel('AQI at time t')
axes[0].set_ylabel('AQI at time t+1')
axes[0].set_title('AQI Persistence (t vs t+1)')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Calculate correlation
corr = df_persist['aqi'].corr(df_persist['aqi_next'])
axes[0].text(0.05, 0.95, f'Correlation: {corr:.3f}',
            transform=axes[0].transAxes, fontsize=12,
            verticalalignment='top', bbox=dict(boxstyle='round',
            facecolor='wheat', alpha=0.5))
```




```

# Distribution of AQI change
axes[1].hist(df_persist['aqi_change'], bins=50, edgecolor='black', alpha=0.7)
axes[1].axvline(0, color='red', linestyle='--', linewidth=2, label='No change')
axes[1].set_xlabel('AQI Change (t+1 - t)')
axes[1].set_ylabel('Frequency')
axes[1].set_title('Distribution of AQI Change')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print(f"\n Key Insight: AQI correlation (t vs t+1) = {corr:.3f}")
print(f"    This suggests persistence is a {'strong' if corr > 0.8 else 'moderate'} baseline.")

```



Key Insight: AQI correlation (t vs t+1) = 0.947
 This suggests persistence is a strong baseline.

```

[52]: # AQI change by hour transition
df_persist['hour_transition'] = df_persist['hour'].astype(str) + ' → ' + \
    df_persist.groupby('station')['hour'].shift(-1).
    ↪astype(str)

transition_stats = df_persist.groupby('hour_transition')['aqi_change'].
    ↪agg(['mean', 'std', 'median'])
print("\nAQI Change by Hour Transition:")
print(transition_stats)

# Plot

```

```


fig, ax = plt.subplots(figsize=(10, 6))
transitions = transition_stats.index[:4] # Get actual transitions
ax.bar(range(len(transitions)), transition_stats.loc[transitions, 'mean'].
    ↪values, alpha=0.7, edgecolor='black')
ax.set_xticks(range(len(transitions)))
ax.set_xticklabels(transitions, rotation=45, ha='right')
ax.set_ylabel('Mean AQI Change')
ax.set_title('Mean AQI Change by Hour Transition')
ax.axhline(0, color='red', linestyle='--', linewidth=1)
ax.grid(True, alpha=0.3, axis='y')
plt.tight_layout()
plt.show()

print("\n Insight: Different hour transitions have different typical changes")
print("    This suggests time-of-day adjusted persistence could be better than_
    ↪simple persistence")

```

AQI Change by Hour Transition:

	mean	std	median
hour_transition			
12 → 18.0	38.26	56.31	26.0
18 → 23.0	-14.60	46.02	-6.0
18 → nan	11.00	44.40	0.0
23 → 6.0	4.27	43.92	0.0
6 → 12.0	-27.95	52.71	-14.0



Insight: Different hour transitions have different typical changes

This suggests time-of-day adjusted persistence could be better than simple persistence

1.13 13. Key Findings Summary

1.13.1 Data Characteristics:

1. **201,664 observations** from 23 stations across 5 cities (2020-2025)
2. **Only 4 measurements per day** (hours: 6, 12, 18, 23)
3. **No missing values** in the dataset
4. **Prediction task:** 5-7 hours ahead (not 1 hour!)

1.13.2 Target Variable (AQI):

- Mean: 265.8, Median: 232.0
- Range: 25-500 (capped at 500)
- Distribution: Right-skewed, dominated by “Severe” and “Moderate” categories
- High persistence: $\text{correlation}(t, t+1) \approx 0.85-0.90$

1.13.3 Strong Predictive Features:

1. **Pollutants** (positive correlation):
 - PM2.5: 0.80
 - PM10: 0.80
 - CO: 0.78
 - NO2: 0.77
 - SO2: 0.76
2. **Weather** (negative correlation):
 - Visibility: -0.86 (strongest predictor!)
 - Temperature: -0.73
 - Wind speed: -0.54

1.13.4 Temporal Patterns:

- **Hour of day:** Hour 18 (6 PM) has highest AQI, Hour 12 (noon) has lowest
- **Season:** Winter » Post-monsoon > Summer > Monsoon
- **Month:** Nov-Jan have highest AQI (winter pollution)
- **Day of week:** Relatively flat (slight weekend effect)

1.13.5 Spatial Patterns:

- **City variation:** Delhi has highest mean AQI
- **Station variation:** Anand Vihar (Delhi) has highest mean AQI (295)
- Clear geographic clustering visible

1.13.6 Implications for Modeling:

1. **Persistence baseline** should perform reasonably well ($\text{corr} \approx 0.85-0.90$)

2. **Lag features** are critical (previous 1-2 measurements)
3. **Visibility** is the single most important weather feature
4. **Temporal features** (hour, season, month) are important
5. **Time-of-day effects** should be modeled (different transitions behave differently)
6. **Station/city effects** exist but can be explored in advanced models